

Loan_Prediction

1 Loan Prediction

```
[ ]: # Import the libraries
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
import seaborn as sns
from sklearn.model_selection import train_test_split
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
```

1:Data Aquisition

```
[ ]: !pip show kaggle
```

```
Name: kaggle
Version: 1.5.16
Summary: Kaggle API
Home-page: https://github.com/Kaggle/kaggle-api
Author: Kaggle
Author-email: support@kaggle.com
License: Apache 2.0
Location: /usr/local/lib/python3.10/dist-packages
Requires: bleach, certifi, python-dateutil, python-slugify, requests, six, tqdm,
urllib3
Required-by:
```

```
[ ]: from google.colab import files
```

```
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving kaggle.json to kaggle.json

```
[ ]: !mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json
```

```
[ ]: # Download the dataset
!kaggle datasets download -d subhamjain/
↳ loan-prediction-based-on-customer-behavior
```

Downloading loan-prediction-based-on-customer-behavior.zip to /content
 78% 4.00M/5.15M [00:00<00:00, 6.02MB/s]
 100% 5.15M/5.15M [00:00<00:00, 5.41MB/s]

```
[ ]: # Unzip the dataset
!unzip loan-prediction-based-on-customer-behavior.zip -d /content/
```

Archive: loan-prediction-based-on-customer-behavior.zip
 inflating: /content/Sample Prediction Dataset.csv
 inflating: /content/Test Data.csv
 inflating: /content/Training Data.csv

2. Data Exploration (EDA) and Preprocessing

```
[ ]: training_df = pd.read_csv("Training Data.csv")
testing_df = pd.read_csv("Test Data.csv")
```

```
[ ]: training_df = training_df.drop(columns=["Id"])
testing_df = testing_df.drop(columns=["ID"])
```

```
[ ]: training_df.head()
```

```
[ ]:
Income  Age  Experience  Married/Single  House_Ownership  Car_Ownership  \
0  1303834   23         3         single         rented         no
1  7574516   40        10         single         rented         no
2  3991815   66         4        married         rented         no
3  6256451   41         2         single         rented        yes
4  5768871   47        11         single         rented         no

      Profession      CITY      STATE  CURRENT_JOB_YRS  \
0  Mechanical_engineer    Rewa  Madhya_Pradesh         3
1  Software_Developer    Parbhani  Maharashtra         9
2   Technical_writer    Alappuzha    Kerala         4
3  Software_Developer  Bhubaneswar    Odisha         2
4   Civil_servant  Tiruchirappalli[10]  Tamil_Nadu         3

CURRENT_HOUSE_YRS  Risk_Flag
0                13         0
1                13         0
2                10         0
```

3	12	1
4	14	1

```
[ ]: testing_df.head()
```

```
[ ]:      Income  Age  Experience  Married/Single  House_Ownership  Car_Ownership  \
0  7393090   59         19         single         rented         no
1  1215004   25          5         single         rented         no
2  8901342   50         12         single         rented         no
3  1944421   49          9        married         rented         yes
4   13429   25         18         single         rented         yes
```

	Profession	CITY	STATE	CURRENT_JOB_YRS	CURRENT_HOUSE_YRS
0	Geologist	Malda	West Bengal	4	13
1	Firefighter	Jalna	Maharashtra	5	10
2	Lawyer	Thane	Maharashtra	9	14
3	Analyst	Latur	Maharashtra	3	12
4	Comedian	Berhampore	West Bengal	13	11

```
[ ]: training_df.columns
```

```
[ ]: Index(['Income', 'Age', 'Experience', 'Married/Single', 'House_Ownership',
          'Car_Ownership', 'Profession', 'CITY', 'STATE', 'CURRENT_JOB_YRS',
          'CURRENT_HOUSE_YRS', 'Risk_Flag'],
          dtype='object')
```

```
[ ]: len(training_df.columns), len(testing_df.columns)
```

```
[ ]: (12, 11)
```

```
[ ]: len(training_df)
```

```
[ ]: 252000
```

```
[ ]: len(testing_df)
```

```
[ ]: 28000
```

```
[ ]: # Check for null values
training_df.isnull().sum()
```

```
[ ]: Income      0
Age            0
Experience      0
Married/Single 0
House_Ownership 0
Car_Ownership  0
Profession      0
```

```
CITY          0
STATE         0
CURRENT_JOB_YRS  0
CURRENT_HOUSE_YRS  0
Risk_Flag     0
dtype: int64
```

```
[ ]: # Check for null values
testing_df.isnull().sum()
```

```
[ ]: Income          0
Age                0
Experience         0
Married/Single    0
House_Ownership   0
Car_Ownership     0
Profession        0
CITY              0
STATE             0
CURRENT_JOB_YRS   0
CURRENT_HOUSE_YRS  0
dtype: int64
```

```
[ ]: # Check for unique values
training_df.index.is_unique
```

```
[ ]: True
```

```
[ ]: testing_df.index.is_unique
```

```
[ ]: True
```

```
[ ]: #Convert the categorical features to numerical features
# List of categorical column names
categorical_columns = ["Married/Single", "House_Ownership", "Car_Ownership",
↳ "Profession", "CITY", "STATE"]

# Create a LabelEncoder instance
label_encoder = LabelEncoder()

# Apply Label Encoding to each categorical column
for column in categorical_columns:
    training_df[column] = label_encoder.fit_transform(training_df[column])
    testing_df[column] = label_encoder.fit_transform(testing_df[column])
```

```
[ ]: training_df.head()
```

```
[ ]:      Income  Age  Experience  Married/Single  House_Ownership  Car_Ownership  \
0  1303834    23         3           1           2           0
1  7574516    40        10           1           2           0
2  3991815    66         4           0           2           0
3  6256451    41         2           1           2           1
4  5768871    47        11           1           2           0

      Profession  CITY  STATE  CURRENT_JOB_YRS  CURRENT_HOUSE_YRS  Risk_Flag
0           33    251     13           3           13           0
1           43    227     14           9           13           0
2           47      8     12           4           10           0
3           43     54     17           2           12           1
4           11    296     22           3           14           1
```

```
[ ]: testing_df.head()
```

```
[ ]:      Income  Age  Experience  Married/Single  House_Ownership  Car_Ownership  \
0  7393090    59         19           1           2           0
1  1215004    25         5           1           2           0
2  8901342    50        12           1           2           0
3  1944421    49         9           0           2           1
4   13429    25        18           1           2           1

      Profession  CITY  STATE  CURRENT_JOB_YRS  CURRENT_HOUSE_YRS
0           26    181     28           4           13
1           24    131     14           5           10
2           30    290     14           9           14
3            1    171     14           3           12
4           12     39     28          13           11
```

Data Visualization

```
[ ]: # Create a function to plot scatter plot between two columns of the dataset
def scatter_plot(x_data, y_data):
    """
    Plots a scatterplot between two columns of the dataset

    Args:
        x_data: column1 of the dataset.
        y_data: column2 of the dataset
    """

    # Create a scatter plot
    plt.scatter(x_data, y_data, color='blue', marker='o', label='Scatter Plot')


    # Add labels and title
    plt.xlabel(x_data.name)
```

```
plt.ylabel(y_data.name)
plt.title(f'{x_data.name} vs {y_data.name}')

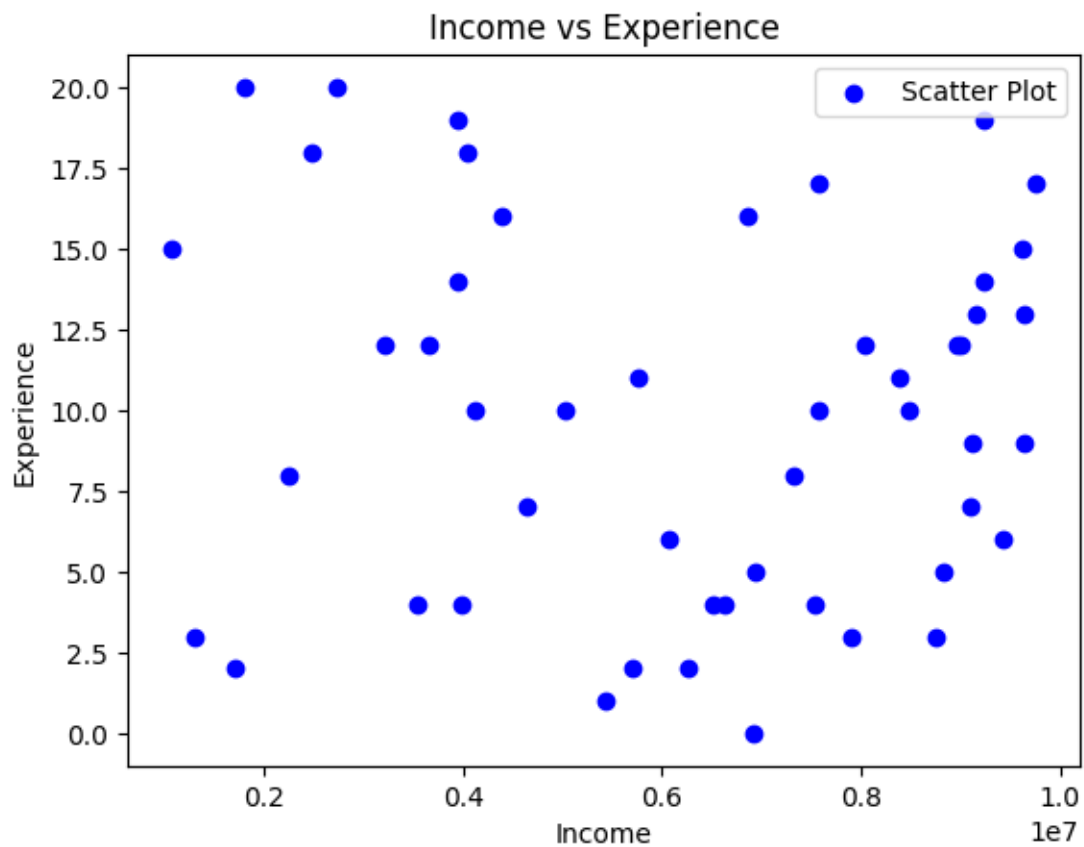
# Add a legend
plt.legend()

# Display the plot
plt.show()
```

```
[ ]: # Plot a scatter plot for the first 50 samples of the training dataset
scatter_plot(training_df["Income"][:50], training_df["Age"][:50])
```




```
[ ]: # Plot a scatter plot for the first 50 samples of the training dataset
scatter_plot(training_df["Income"][:50], training_df["Experience"][:50])
```




```
[ ]: # Plot Histogram for age column
sns.histplot(training_df.Age)
plt.xlabel("Age")
plt.title("Age Count")
```

```
[ ]: Text(0.5, 1.0, 'Age Count')
```




```
[ ]: # Plot Box plot on Income column
sns.boxplot(training_df.Income)
```

```
[ ]: <Axes: ylabel='Income'>
```

```
[ ]: # Plot Heat map
plt.figure(figsize=(10,8))
sns.heatmap(training_df.corr(),annot=True,cmap='cividis')

plt.show()
```



3: Splitting into training and testing set

```
[ ]: x = training_df.drop("Risk_Flag", axis=1)
     y = training_df["Risk_Flag"]
```

```
[ ]: # Split into train & test set
     x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

4. Modelling

```
[ ]: # Put models in a dictionary
     models = {"Logistic Regression": LogisticRegression(),
               "KNN": KNeighborsClassifier(),
               "Random Forest": RandomForestClassifier()}

     # Create a function to fit and score models
```


```
def fit_and_score(models, X_train, X_test, y_train, y_test):
    """
    Fits and evaluates given machine learning models.
    models : a dict of different Scikit-Learn machine learning models
    X_train : training data (no labels)
    X_test : testing data (no labels)
    y_train : training labels
    y_test : test labels
    """
    # Set random seed
    np.random.seed(42)
    # Make a dictionary to keep model scores
    model_scores = {}
    # Loop through models
    for name, model in models.items():
        # Fit the model to the data
        model.fit(X_train, y_train)
        # Evaluate the model and append its score to model_scores
        model_scores[name] = model.score(X_test, y_test)
    return model_scores
```

```
[ ]: model_scores = fit_and_score(models=models,
                                   X_train=x_train,
                                   X_test=x_test,
                                   y_train=y_train,
                                   y_test=y_test)

model_scores
```

```
[ ]: {'Logistic Regression': 0.8767857142857143,
      'KNN': 0.888234126984127,
      'Random Forest': 0.8980952380952381}
```

```
[ ]: # Comparing the models
model_compare = pd.DataFrame(model_scores, index=["accuracy"])
model_compare.T.plot.bar();
```



Hyperparameter Tuning

```
[ ]: train_scores = []
test_scores = []

# List of different values of n_neighbours
neighbors = range(1, 15)

# Initialize KNN
knn = KNeighborsClassifier()

for i in neighbors:
    knn.set_params(n_neighbors=i)
    knn.fit(x_train, y_train) # Fitting the algorithm
    train_scores.append(knn.score(x_train, y_train)) # Updating train_scores
```

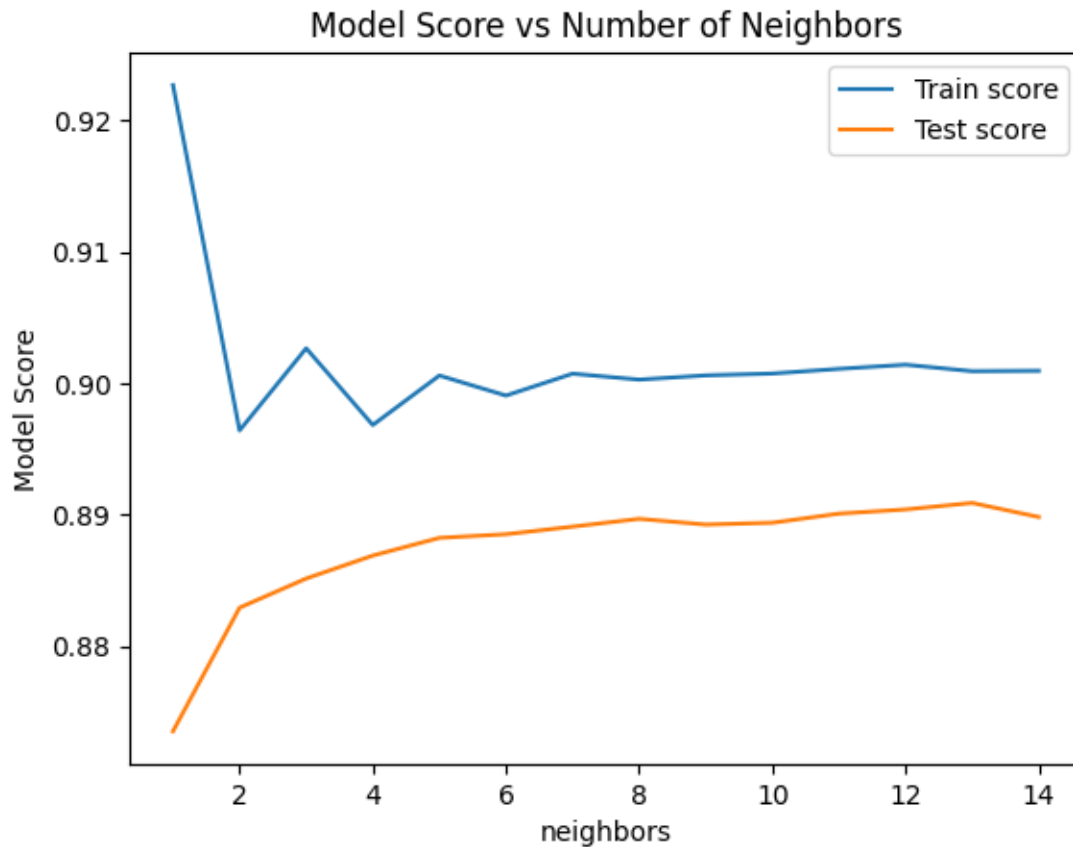
```
test_scores.append(knn.score(x_test, y_test)) # updating test_scores
```

```
[ ]: train_scores, test_scores
```

```
[ ]: ([0.9227083333333334,  
      0.8964136904761905,  
      0.9026686507936508,  
      0.8968253968253969,  
      0.9005952380952381,  
      0.8990674603174603,  
      0.900734126984127,  
      0.9002827380952381,  
      0.9006051587301588,  
      0.9007440476190476,  
      0.9011011904761905,  
      0.9014186507936508,  
      0.900922619047619,  
      0.900952380952381],  
      [0.8735119047619048,  
      0.8829166666666667,  
      0.8851388888888889,  
      0.8868849206349206,  
      0.888234126984127,  
      0.8885119047619048,  
      0.8890873015873015,  
      0.8896825396825396,  
      0.8892460317460318,  
      0.8893849206349206,  
      0.890079365079365,  
      0.8903968253968254,  
      0.8908928571428572,  
      0.8898214285714285])
```

```
[ ]: # Plot the model scores  
plt.plot(neighbors, train_scores, label="Train score")  
plt.plot(neighbors, test_scores, label="Test score")  
plt.xlabel("neighbors")  
plt.ylabel("Model Score")  
plt.title("Model Score vs Number of Neighbors")  
plt.legend()  
print(f"Maximum model score: {max(test_scores)*100:.2f}%")
```

Maximum model score: 89.09%



```
[ ]: # Tuning with RandomizedSearchCV

# Define the hyperparameter search space for LogisticRegression
log_param_dist = {
    'C': np.logspace(-5, 5, 20), # Regularization strength
    'penalty': ['l1', 'l2'], # Type of regularization
    'solver': ['saga'] # Solver algorithm
}

[ ]: # Tune LogisticRegression model
# Setup random seed
np.random.seed(42)

# Setup random hyperparameter search for logistic regression
rs_log_reg = RandomizedSearchCV(LogisticRegression(),
                                param_distributions=log_param_dist,
                                cv=5,
                                n_iter=20,
                                verbose=True)
```



```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```



```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350:
ConvergenceWarning: The max_iter was reached which means the coef_ did not
converge

```

```

[ ]: RandomizedSearchCV(cv=5, estimator=LogisticRegression(), n_iter=20,
                        param_distributions={'C': array([1.00000000e-05,
3.35981829e-05, 1.12883789e-04, 3.79269019e-04,
1.27427499e-03, 4.28133240e-03, 1.43844989e-02, 4.83293024e-02,
1.62377674e-01, 5.45559478e-01, 1.83298071e+00, 6.15848211e+00,
2.06913808e+01, 6.95192796e+01, 2.33572147e+02, 7.84759970e+02,
2.63665090e+03, 8.85866790e+03, 2.97635144e+04, 1.00000000e+05]),
                        'penalty': ['l1', 'l2'],
                        'solver': ['saga']},
                        verbose=True)

```

```

[ ]: rs_log_reg.best_params_

```

```

[ ]: {'solver': 'saga', 'penalty': 'l2', 'C': 0.5455594781168515}

```

```

[ ]: # Evaluate the model
rs_log_reg.score(x_test, y_test)

```

```

[ ]: 0.8767857142857143

```

```

[ ]:

```

5 Model Selection

```
[ ]: print(f"KNN Model (Tuned): {max(test_scores)*100:.2f}%")
      print(f"Logistic Regression (RandomizedSearchCV): {(rs_log_reg.score(x_test, y_test))*100:.2f}%")
```

KNN Model (Tuned): 89.09%

Logistic Regression (RandomizedSearchCV): 87.68%

```
[ ]: model_scores
```

```
[ ]: {'Logistic Regression': 0.8767857142857143,
      'KNN': 0.888234126984127,
      'Random Forest': 0.8980952380952381}
```

```
[ ]: rf_model = RandomForestClassifier()
      rf_model.fit(x_train, y_train)
      rf_model.predict(x_test)
      rf_model.score(x_test, y_test)
```

```
[ ]: 0.8979166666666667
```

```
[ ]: # Since Random Forest has highest performarnce we will select Random Forest for predictions
```

```
sample_data = {"Income": 13023043,
                "Age": 33,
                "Experience": 10,
                "Married/Single": 1,
                "House_Ownership": 3,
                "Car_Ownership": 5,
                "Profession": 433,
                "CITY": 220,
                "STATE": 15,
                "CURRENT_JOB_YRS": 11,
                "CURRENT_HOUSE_YRS": 16}

sample_df = pd.DataFrame(sample_data, index=[0])
sample_df
prediction = rf_model.predict(sample_df)
if prediction[0] == 0:
    print(f"Prediction: {prediction[0]} \n No risk")
else:
    print(f"Prediction: {prediction[0]} \n High Risk")
```

Prediction: 0

No risk

```
[ ]:
```